

Economic-Aware and Sovereignty-Constrained Routing for Enterprise LLM Gateways

Draft v0.1 — single-provider (OpenAI) evaluation. This is a first complete write-up built from

real meas

Abstract

Enterprises are adopting large language models (LLMs) faster than they can govern them: they cannot

say which
model,

1. Introduction

LLMs are now embedded in HR assistants, document RAG, developer tooling, and analytical agents. Each

call has a

1. **Cost opacity.** Teams cannot attribute spend per application/team or decide when a cheaper model

would suff

Existing gateways and FinOps tools mostly *measure and report*. Our thesis is that a control plane at

the **gatew**

Contributions.

- A cons

2. Background and Related Work

LLM gateways. Envoy AI Gateway and LiteLLM provide a unified data path to multiple providers

Model routing and cascades. Prior work routes between models by difficulty or confidence, or

FinOps for AI. Cloud FinOps practice emphasizes attribution and budgets

Serving economics. Self-hosting via vLLM/TGI trades fixed GPU/ops cost for marginal token cost

3. Problem Formulation and Algorithm

A request r carries metadata (team, namespace, sensitivity) and an estimated token footprint. Given

Hard constraint. Reject any m that violates P : its zone is forbidden; or an allow-list is set

Soft objective. Among C meeting a minimum quality $q_{\min}$ (relaxed only when the budget is

$$\arg\min_{m \in C}; \quad \alpha \cdot \widehat{c}(m) \cdot (1 + \epsilon \cdot \rho) + \beta \cdot \ell_q(m) + \gamma \cdot \ell_\lambda(m) + \delta \cdot \sigma(m)$$

where $\widehat{c}(m) = \text{cost}(m) / \text{cost}(\text{premium})$ is normalized cost, $\rho = \text{used} / \text{total}$

Pseudocode and the exact normalization are in `methodology.md`; the implementation is

4. System Design

Requests flow client → **Envoy AI Gateway** → provider (managed API or self-hosted). The **AI**

Sovereign
budgets,

5. Methodology

Testbed. A Go harness drives the operator's **actual** engines and issues **real** LLM calls. The

provider in

Workloads (W1-W4). HR chatbot (short, sensitive), RAG docs (long context, quality-sensitive), Dev

assistant (

Baselines. B1 premium-static, B2 round-robin, B3 least-cost, B4 static namespace policy, B5 budget

hard-block

Metrics. Cost (total, per request, per token, per team via the operator cost engine, % savings);

quality (LL
the

Reproducibility. Temperature 0; all responses and judge scores cached (``results/cache.json``) so

re-runs are

6. Results

RQ1 — Cost (Figure 2, Table 1)

Relative to

RQ2 — Quality preservation (Figure 3, Table 2)

B6-ours
97.5%),

RQ3 — Latency and overhead (Figure 4, Table 3)

The **routing**

RQ4 — Sovereignty (Figure 5, Table 4)

RQ5 — Budget-aware degradation (Figure 6, Table 5)

Under a b
availabili

RQ6 — Managed vs self-hosted break-even (Figure 7) — **modeled**

Using rea
managed

Ablation (Figure 8, Table 7)

Removing
term

7. Discussion

Economic-aware routing is most valuable when a workload mixes easy and hard requests: the router sends

cheap req

8. Threats to Validity

Summarized here; full version in `threats\_to\_validity.md`. **Internal:** API/latency variability

(mitigated

9. Conclusion and Future Work

A gateway-level, economic-aware, sovereignty-constrained, budget-aware control plane can cut LLM cost

substantia

Artifact / Reproducibility

Operator + `experimentation/` (datasets, harness, scripts, results, figures, cached responses).

Reproduce

Appendix: Figures

Figure 2 — Total cost by routing strategy

0.040
0.035
0.030
0.025
0.020
0.015
0.010
0.005
0.000

Total cost (EUR)

Figure 3 — Quality vs cost

0.90
0.89
0.88
0.87
0.86

Mean quality (normalized 0-1)

Figure 4 — Latency p95 and mean $\pm$ 95% CI

3000 -
2500 -
2000 -
1500 -
1000 -
500 -
0 -

Latency (ms)

Figure 5 — Sovereignty: cost & violations




Figure 6 — Budget policy: availability vs overrun



Figure 7 — Managed vs self-hosted break-even (modeled)




Figure 8 — Ablation of scoring terms

